

Kite Logik

The Governance Control Plane for AI Agents

Deterministic, auditable policy enforcement
at the infrastructure level - not the prompt level.

TETHER

OPA Policy Engine

SANDBOX

MicroVM Isolation

ANCHOR

HITL + Audit

WHAT IS KITE LOGIK

Kite Logik is the **governance control plane for AI agents**. It governs what agents can do, what they can spawn, what data they can access, and what resources they can consume - enforced at the infrastructure level, not the prompt level.

"A prompt-based guardrail is a suggestion. Kite Logik is a lock."

The core distinction: every competing tool relies on the language model cooperating with its instructions. Kite Logik removes that dependency. Policy decisions are made by OPA - a deterministic, testable policy engine - and enforced by the sandbox and credential broker. The model cannot argue past a deny.

THE PROBLEM IT SOLVES

As agentic AI moves from demos to production, the question shifts from *"can this agent complete the task?"* to *"should this agent be permitted to take this action?"* Today's tooling answers only the first question.

Approach	Example Tools	The Gap
Prompt-level guards	Guardrails AI, NeMo	Suggestions - the model can be prompted to ignore them
Evaluation / testing	Promptfoo, Patronus AI	Pre-deployment only, not runtime enforcement
ML classifiers	LlamaGuard	Probabilistic, non-auditable, no policy versioning
Cloud-locked guardrails	AWS Bedrock, Azure AI	Vendor lock-in, no Rego, no HITL, no sandbox
Infrastructure enforcement	Kite Logik	The gap Kite Logik fills

Market timing is acute. The EU AI Act's Article 9 risk-management obligations came into force in 2025, with high-risk AI system requirements applying from 2026. SOC 2 Type II for AI agent deployments is becoming a standard procurement requirement, and no existing tooling generates the required audit evidence automatically.

HOW IT WORKS - THREE LAYERS

Every governance event - tool calls, agent spawns, delegation requests, proposed plans, resource consumption, and data access - flows through the same three-layer pipeline:

1 TETHER - Policy Engine	2 SANDBOX - Isolation	3 ANCHOR - Oversight
OPA/Rego evaluates every event before execution. Denies if policy explicitly forbids the action.	Docker containers (OS) enforce risk to sandbox. Escapes with the system default privilege.	Control plane with the system default privilege.

WHAT KITE LOGIK GOVERNS

All six categories of agent behavior are evaluated through a single, unified policy pipeline:

Governance Event	What Is Evaluated	Example Policy
Tool Calls	Every tool invocation, before execution	Block file writes outside /tmp
Agent Spawn	Agent creation with requested capabilities	Max delegation depth is 2
Delegation	Agent-to-agent task handoff	Child scopes must be subset of parent's
Plans	Proposed action sequence, before any step runs	Deny plans containing blocked tools
Resource Budgets	Token spend, API calls, compute time	Deny if session budget exhausted
Data Access	Classification-based flow control	Confidential data stays in primary session

On Human-in-the-Loop (HITL): HITL is the exception, not the default. The standard path - OPA evaluates, infrastructure enforces - adds zero human latency. HITL fires only when OPA policy explicitly sets `requires_hitl := true`, reserved for narrow, high-stakes scenarios: wire transfers above threshold, restricted data classification access, or actions flagged as security-critical.

FOR DEVELOPERS - OSS QUICK START

The open-source release provides the **full governance pipeline** - not a feature-gated trial. Add Kite Logik to any agent in three lines of code:

```
from kitelogik import governed, PolicyGate, OPAClient, SessionContext
gate = PolicyGate(opa_client=OPAClient())
ctx = SessionContext(
    session_id="s1", user_role="support",
    session_scopes=["read_customer", "approve_refund_under_100"]
)
@governed(gate=gate, context=ctx)
async def approve_refund(customer_id: str, amount: float) -> str:
    return payment_api.refund(customer_id, amount)
# approve_refund('cust_123', 50.0) -> OPA allows, runs normally
# approve_refund('cust_123', 500.0) -> OPA denies, raises GovernanceError
```

Included in OSS	
OPA/Rego policy engine, deny-by-default	Immutable audit trail with integrity hashing
Tool call, lifecycle, plan, budget, and data governance	OpenTelemetry tracing
Docker sandbox with network isolation	Real-time WebSocket dashboard
Session-scoped credentials with delegation	Starter policy library (financial, security, ops)
HITL escalation queue	OpenAI and LangChain framework adapters

POLICY-AS-CODE WITH OPA/REGO

Kite Logik uses **Open Policy Agent (OPA)** as its policy engine - the same tool security teams already use for Kubernetes admission control. Policies are written in Rego, version-controlled, testable offline, and deployed without touching application code.

```
package kitelogik.financial
default allow := false
# Allow read-only operations for all roles
allow if { input.action == "read" }
# Require HITL for high-value transactions
requires_hitl if {
    input.action == "approve_refund"
    input.args.amount > 1000
}
# Block amounts above hard limit entirely
deny if {
    input.action == "approve_refund"
    input.args.amount > 10000
}
```

The same policy evaluates in milliseconds, produces a structured audit record, and can be unit-tested with OPA's built-in test runner - entirely independent of the AI model.

FOR ENTERPRISES - ORGANIZATIONAL-SCALE GOVERNANCE

The enterprise tier extends Kite Logik from single-agent governance to fleet-wide, multi-tenant, compliance-ready infrastructure for regulated environments.

Enterprise Feature	Capability	Relevant For
Firecracker MicroVM	Hardware-enforced kernel isolation, sub-125ms boot	Fintech, GovCon, HIPAA
Postgres Backends	HA durable storage, connection pooling, read replicas	Production multi-node
SSO + RBAC	SAML/OIDC; Admin / Policy Author / Operator / Viewer	Team and org management
Compliance Export Packs	One-click SOC 2, HIPAA, FedRAMP evidence (PDF/CSV/JSON)	ISO and procurement
Policy Intelligence	Decision heatmaps, anomaly detection, HITL SLA tracking	SecOps and risk teams
Agent Fleet Management	Real-time inventory, remote kill switch, canary deploys	Platform engineering
Multi-Tenant Isolation	Namespace-isolated OPA bundles per team, central registry	SaaS multi-tenant
Managed SIEM Connectors	Splunk, Datadog, Elastic with retry and backpressure	Enterprise SecOps
Policy Simulation	Replay historical sessions against new policy drafts	Risk-free policy changes
Cross-Agent Governance	Org-wide budgets, concurrent access limits	Fleet coordination

DEPLOYMENT MODES

Mode	Who Uses It	How It Operates
Embedded SDK	Individual developers	@governed decorator wraps tool functions in-process. Zero network hop. Governance
Governance Gateway	Platform / DevSecOps teams	FastAPI service that agents call tools through. Centralized enforcement, fleet-wide ob
Both	Enterprise deployments	SDK for in-process speed; Gateway for centralized audit, HITL management, and com

Tier	Price	Target	Key Differentiator
OSS	Free	Developers, small teams	Full runtime, SQLite, Apache 2.0
Enterprise	\$2,500-\$8,000/mo	Platform teams, regulated orgs	Postgres, Firecracker, SSO/RBAC, compliance exports
Enterprise Plus	Custom	FedRAMP, healthcare, large-scale	All above + SLA support, dedicated CSM, security review

One developer-month at loaded cost (\$25-40k) equals approximately one year of Enterprise. Building the equivalent in-house - OPA evaluation, HITL queue, credential broker, sandbox lifecycle, memory provenance, OpenTelemetry - represents 6-12 months of engineering.

COMPETITIVE POSITIONING

Kite Logik is **not** a replacement for Promptfoo, Guardrails AI, or NeMo Guardrails. Each addresses a different layer of the AI safety stack, and Kite Logik is designed to be complementary - not competitive - with all of them.

Tool	Primary Function	When It Runs	Relationship to Kite Logik
Promptfoo	LLM evaluation and red teaming	Pre-deployment	Complementary - test prompts before shipping
Guardrails AI	LLM output content validation	Post-generation	Complementary - validate content quality
NeMo Guardrails	Conversational flow control	Dialog runtime	Complementary - manage conversation safety
Kite Logik	Agent action governance	Tool call runtime	The enforcement layer the others don't provide

"Use Promptfoo to test your prompts. Use Guardrails AI to validate your outputs. Use Kite Logik to govern what your agents actually execute."

UNIQUE CAPABILITIES - NO COMPETITOR HAS THESE

Capability	Why It Matters
Infrastructure-level enforcement	Violations are structurally impossible - not reliant on model cooperation
OPA/Rego policy language	Security teams already know it from Kubernetes; zero new language to learn
Sandbox isolation (Docker/MicroVM)	Network deny-by-default; agent cannot reach blocked APIs regardless of instructions
Agent lifecycle governance	Governs spawn, delegation, and plans - not just individual tool calls
Plan-before-execute gate	Evaluates intent across a proposed sequence, before any step runs
Session-scoped credentials	Per-session tokens auto-revoked on termination; no credential leakage
Memory provenance tracking	Every memory write carries source, trust tier, session ID, and timestamp
Async HITL escalation	Human review without blocking real-time execution

REGULATORY ALIGNMENT

Regulation / Standard	Relevant Requirement	How Kite Logik Helps
EU AI Act - Art. 9	Risk management for high-risk AI	Policy-gated tool calls, audit trail, HITL for critical actions
EU AI Act - Art. 15	Accuracy, robustness, cybersecurity (from 2026)	Sandbox isolation, credential scoping, memory provenance
SOC 2 Type II	Evidence of access controls, audit logging	Immutable audit log, session credentials, compliance exports
HIPAA	PHI access controls, audit trails	Data classification policies, memory trust tiers, audit chain
FedRAMP	Continuous monitoring, least privilege	Deny-by-default network, OPA policy versioning, kill switch

NEXT STEPS

For Developers

Install the OSS package and add governance to your first agent in under five minutes.

```
pip install kitelogik
```

The quickstart demo runs five scenarios - ALLOW, BLOCK, HITL, indirect prompt injection defence, and multi-agent delegation - from a single command with no external dependencies.

Start at github.com/kitelogik/kitelogik or read the documentation at docs.kitelogik.io.

For Enterprise Teams

We are currently accepting **three design partners** for the enterprise tier - one each from fintech, healthtech, and government/contractor verticals.

Design partners receive:

- A free 12-month enterprise licence
- 4 hours/month architecture consultation with the founder
- Direct input on the compliance evidence roadmap

Contact us at enterprise@kitelogik.io to discuss whether your deployment context is a fit.

THE ONE-SENTENCE PITCH

**You use OPA for your Kubernetes pods.
Kite Logik lets you use it for your AI agents.**

Kite Logik is released under the Business Source Licence 1.1 (core runtime) and Apache 2.0 (peripheral libraries). On-premise and private-cloud packaging available. No usage data leaves your environment.

kitelogik.io · github.com/kitelogik · enterprise@kitelogik.io