

KITE LOGIK

The Governance Control Plane for AI Agents

Product Overview

kitelogik.com

April 2026

What Kite Logik Is

Kite Logik is the governance control plane for AI agents. It governs what agents can do, what they can spawn, what they can access, and what resources they can consume — enforced at the infrastructure level, not the prompt level.

Other tools test prompts. Other tools validate LLM outputs. Kite Logik governs the agent itself.

| *A prompt-based guardrail is a suggestion. Kite Logik is a lock.*

The key distinction: prompt-level guardrails rely on the model cooperating. Kite Logik does not. Policies are evaluated by OPA, enforced by the credential broker and policy gate. The model cannot override a deny.

What It Governs

Kite Logik governs six categories of agent behavior through a single policy pipeline:

Governance Event	What's Evaluated	Example
Tool calls	Every tool invocation, before execution	"Block file writes outside /tmp"
Agent spawn	Agent creation with requested capabilities	"Max delegation depth is 2"
Delegation	Agent-to-agent task handoff	"Child scopes must be subset of parent"
Plans	Proposed action sequence, before any step runs	"Deny plans containing blocked tools"
Resource budgets	Token spend, API calls, compute time	"Deny if session budget is exhausted"
Data access	Classification-based flow control	"Confidential data stays in primary session"

Every governance event flows through the same pipeline:

Governance Event → Credential Check → OPA Evaluation → ALLOW / DENY / HITL Escalation

HITL is the exception, not the default. It fires only when OPA policy explicitly requires it — for wire transfers, restricted data access, or security-critical actions. The default governance path adds zero human latency.

Architecture: Two Layers

1. Tether (Policy Engine) — OPA/Rego evaluation of every governance event. Deterministic, testable, version-controlled. The same policy language security teams already use for Kubernetes admission control. Supports both a standalone OPA server and Regorus (an in-process Rust-based Rego engine) — zero external infrastructure required to get started.

2. Anchor (Oversight) — Session-scoped credentials with automatic revocation. Async HITL queue for high-stakes decisions. Immutable audit trail with integrity hashing. Full OpenTelemetry tracing across every policy decision.

Enterprise Edition adds a third layer: Sandbox (Isolation) — Docker containers and Firecracker MicroVMs with network isolation, resource limits, and hardware-enforced kernel isolation. One container per agent session.

How It's Deployed

Mode	Who Uses It	How
Embedded SDK	Individual developers and teams	@governed decorator, GovernedToolbox, 11 framework adapters. Runs in-process. Zero network hop.
Governance Gateway	Platform teams (Enterprise)	FastAPI HTTP API for centralized enforcement, fleet-wide observability, and compliance reporting.

Most developers start with the Embedded SDK — it covers the full governance pipeline with no external services beyond an optional OPA server. Enterprise deployments add the Gateway for centralized policy enforcement and fleet management.

How Developers Use the Open Source

The open-source release gives you the **full governance pipeline** for your agents — not a feature-gated trial.

3-Line Integration

```
from kitelogik import governed, PolicyGate, RegorusClient, SessionContext

gate = PolicyGate(opa_client=RegorusClient()) # in-process, no server needed
ctx = SessionContext(session_id="s1", user_role="support",
                     session_scopes=["read_customer",
                                     "approve_refund_under_100"])

@governed(gate=gate, context=ctx)
async def approve_refund(customer_id: str, amount: float) -> str:
    return payment_api.refund(customer_id, amount)

# approve_refund("cust_123", 50.0)    → OPA allows, runs normally
# approve_refund("cust_123", 500.0)   → OPA denies, raises GovernanceError
```

Framework Adapters

11 framework adapters ship out of the box. All share the same governance pipeline — intercept tool calls, route through policy, return standard results:

- OpenAI, Anthropic, LangChain, LangGraph, CrewAI
- OpenAI Agents SDK, Google ADK, PydanticAI, LlamaIndex, Semantic Kernel, Haystack

KITE LOGIK

What's Included in OSS

Everything you need to govern AI agents — structurally, not probabilistically.

Policy Engine

- ▶ OPA/Rego policy engine with deny-by-default enforcement
- ▶ Regorus in-process Rego engine — no OPA server or Docker required
- ▶ YAML policy frontend (kitelogik compile) — write policies without learning Rego
- ▶ 2-tier policy hierarchy (global + project)

Governance Scope

- ▶ Tool call, agent lifecycle, plan, budget, and data classification governance
- ▶ Session-scoped credentials with delegation and scope narrowing

Human Oversight

- ▶ HITL queue for high-stakes decisions (code-level API)

Audit & Observability

- ▶ Immutable audit trail with SQL trigger enforcement
- ▶ OpenTelemetry tracing across every policy decision

Agent Memory

- ▶ Agent memory with trust tiers and provenance tracking

Developer Tools

- ▶ Starter policy library with ready-to-use templates
- ▶ 11 framework adapters
- ▶ CLI for policy compilation, validation, and compliance audit
- ▶ Compliance CLI with OWASP ASI mapping

What's not in OSS: Organizational-scale features — see Enterprise section below.

How Enterprises Use It

Enterprise adds governance at **organizational scale** with Kite Logik Enterprise:

Infrastructure & Isolation

- **Firecracker MicroVM** — Hardware-enforced kernel isolation. Sub-125ms boot. Required for fintech and government.
- **Docker sandbox** — Network isolation, resource limits, read-only rootfs. One container per agent session.

Data & Storage

- **PostgreSQL backends** — HA, connection pooling, read replicas for HITL queue, credentials, memory, and audit.

Access & Identity

- **SSO + RBAC** — SAML/OIDC single sign-on. Admin, Policy Author, Operator, Viewer roles scoped to tenant and team.

Observability & Compliance

- **Real-time governance dashboard** — Live event feed, HITL approve/deny UI, memory browser, fleet view, audit log, and trace viewer.
- **Prometheus metrics** — Decision counters, HITL queue gauges, gate latency histograms with Grafana dashboards.
- **Managed SIEM connectors** — Splunk HEC, Datadog, Elastic with retry and backpressure.
- **Compliance export packs** — One-click SOC 2, HIPAA, FedRAMP audit evidence in PDF, CSV, and JSON.
- **Policy intelligence dashboard** — Decision heatmaps, anomaly detection, session risk scoring, HITL SLA tracking.

Fleet & Operations

- **Governance Gateway** — Centralized HTTP API for policy enforcement across multiple agents. REST endpoints for HITL approval.
- **Multi-agent Orchestrator** — Spawn and coordinate multiple governed agent sessions with delegation chains.
- **Agent fleet management** — Real-time inventory, remote kill switch, canary policy deployments.
- **Multi-tenant policy isolation** — Namespace-isolated OPA bundles per team. Central registry with inheritance.
- **Cross-agent governance** — Org-wide budgets, concurrent access limits, coordinated deny policies.
- **Policy simulation** — Replay historical sessions against proposed policy changes before deploying.
- **Governance marketplace** — Community and enterprise compliance packs.

The enterprise model: OSS proves the value on a single agent. Enterprise scales it across the organization.

For enterprise licensing: licensing@kitelogik.com

Competitive Positioning

Kite Logik is **not competing** with Promptfoo, Guardrails AI, or NeMo Guardrails. These tools are complementary:

Tool	What It Does	Category
Promptfoo	Tests whether your LLM is safe	Pre-deployment evaluation
Guardrails AI	Validates LLM outputs	Content quality
NeMo Guardrails	Controls conversation flow	Dialog safety
Kite Logik	Enforces what agents can do	Runtime governance

Use Promptfoo to test your prompts. Use Guardrails AI to validate your outputs. Use Kite Logik to govern what your agents actually execute.

What No Competitor Has

OPA/Rego policy-as-code, HITL escalation, session-scoped credentials with delegation, agent lifecycle governance, plan-before-execute validation, delegation chain enforcement, memory provenance tracking, YAML policy frontend, and 11 framework adapters out of the box.

The Core Thesis

As agents move from demos to production, the question shifts from “*Can this agent do the task?*” to “*Should this agent be allowed to do the task?*”

Kite Logik answers the second question at the infrastructure level, where it cannot be prompt-injected away.

Kite Logik — The Governance Control Plane for AI Agents

kitelogik.com · [pip install kitelogik](#) · licensing@kitelogik.com